

Can Insider Trading be Systematically Detected?

GAN Overview

Context

- Large body of literature attempting to identify insider trading with limited success
- Prior attempts use regressions / linear approaches
- New machine learning approaches untested

Aim vs. Challenges

Aim

- Develop a ML algorithm that can systematically detect insider trading

Challenges

- Finance literature has been unable to identify any observable features of insider trading
- Prosecuted IT cases are sparse → if you train a ML algorithm on the IT data then there isn't often any out of sample data to validate algorithm
- ML algorithms typically biased towards majority class
- Prosecuted IT cases represent a minimum bound → A successful model would identify all IT periods + some

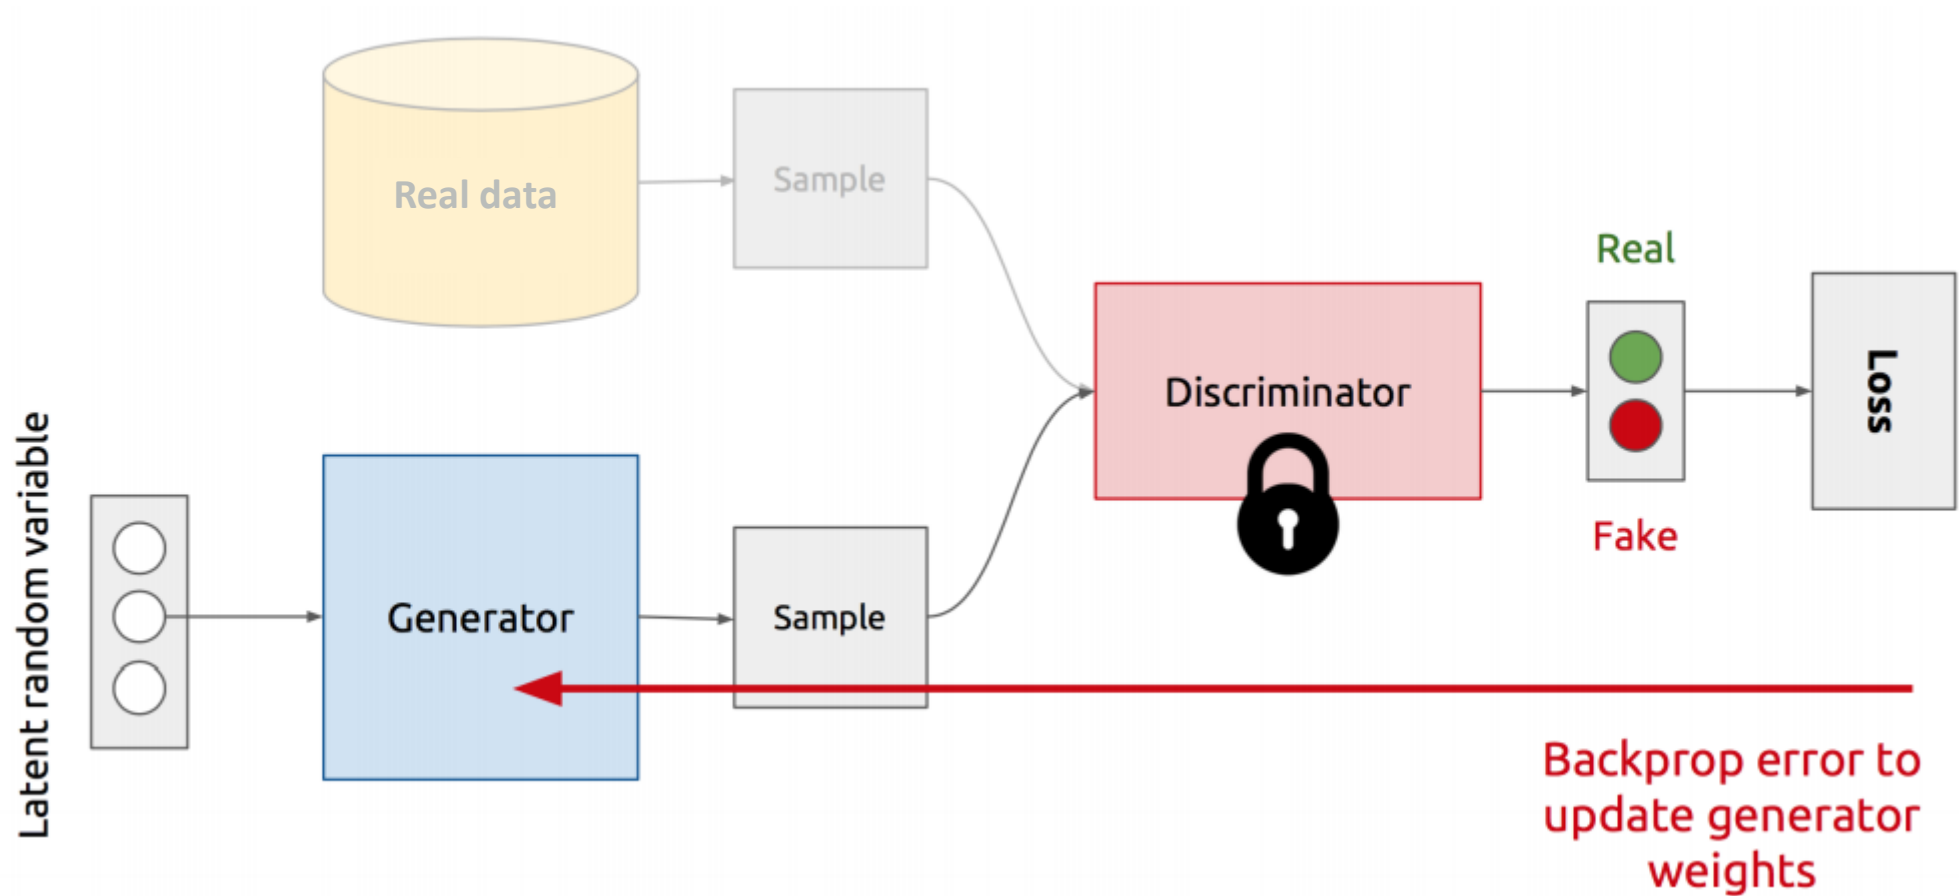
How can we train a model to detect insider trading without the model ever observing insider trading?



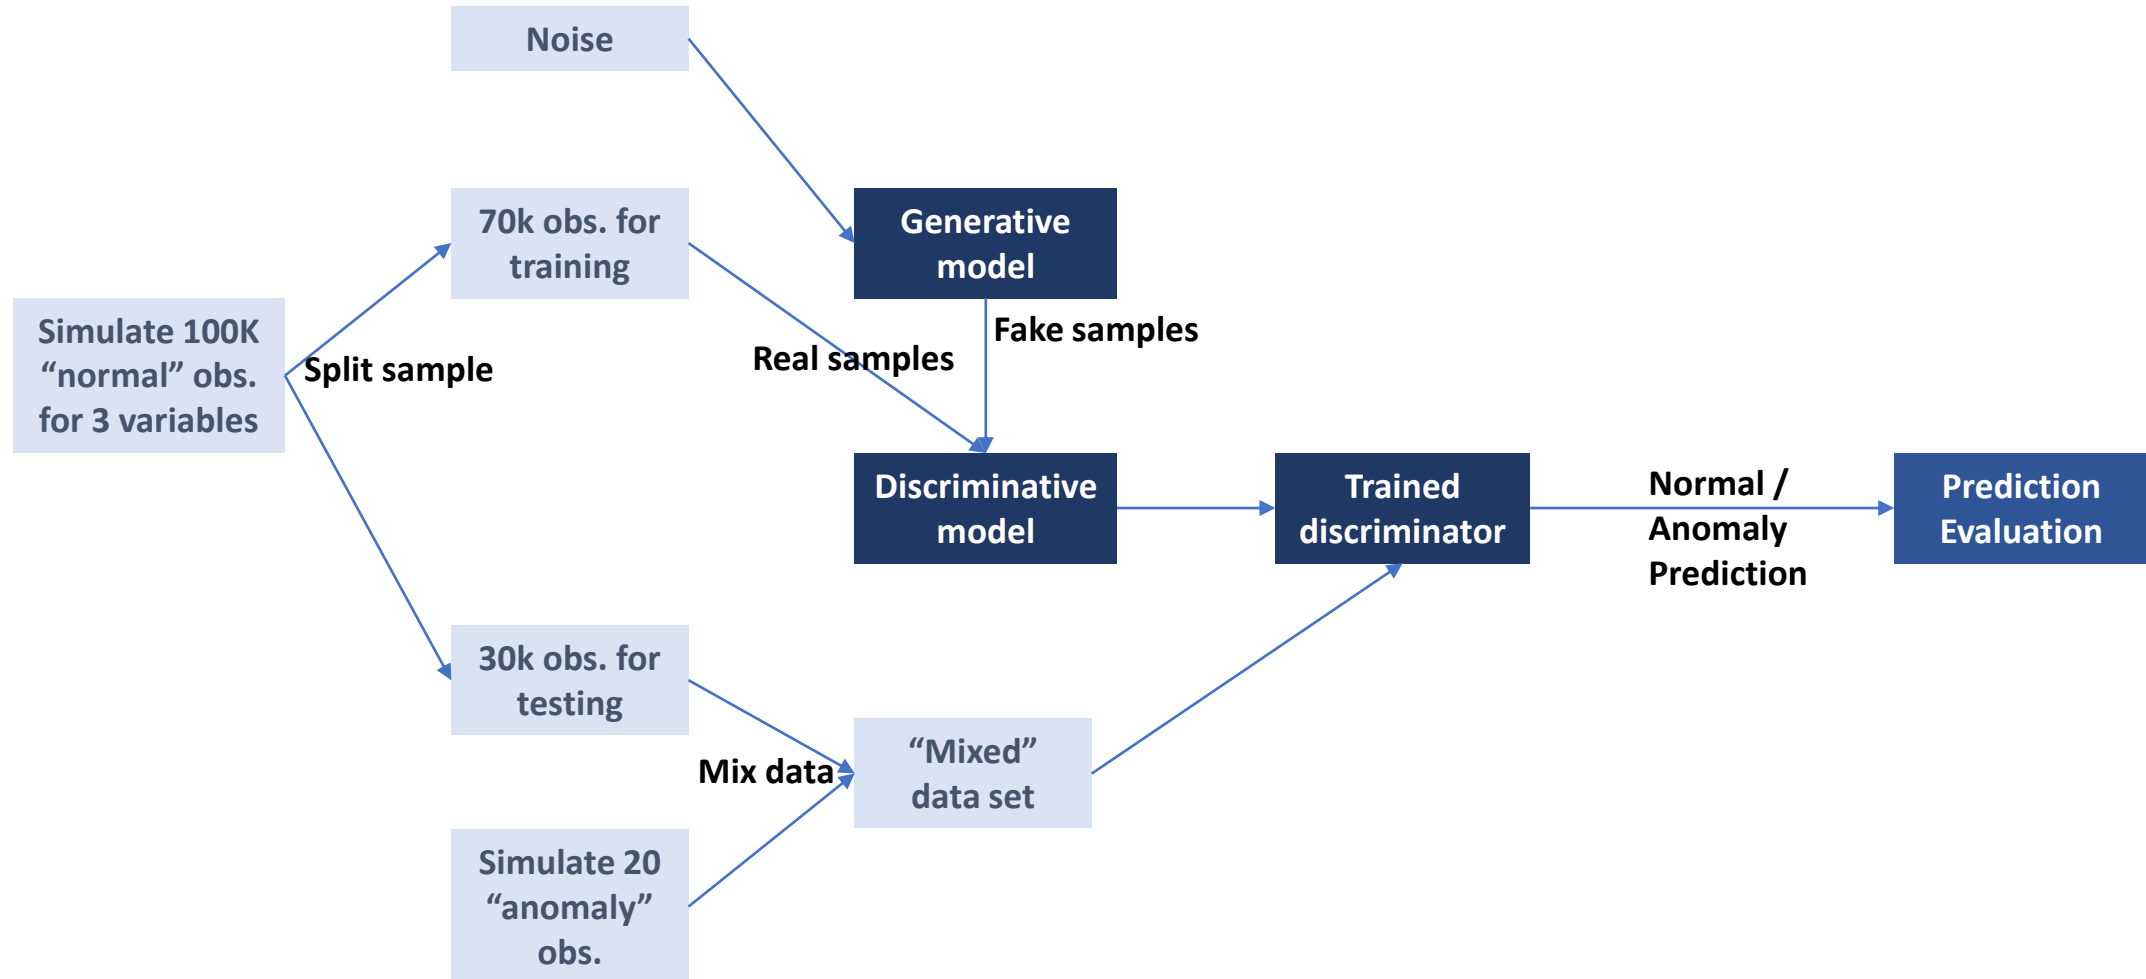
GANs provide a suitable framework

- We can train GANs to learn a representation of ordinary market data
- The trained model can then flag abnormal data
- We want to observe whether abnormal data flagged by the trained model matches prosecuted cases of IT

Standard GAN framework

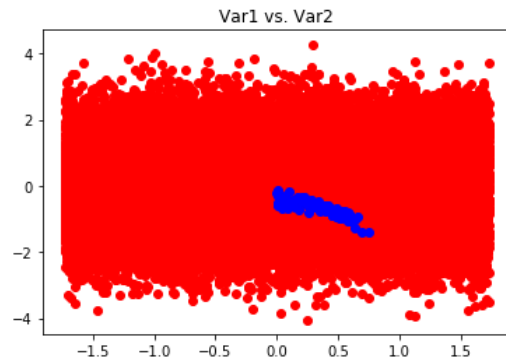


An example workflow with simulated data




Example training progress


Epoch 1



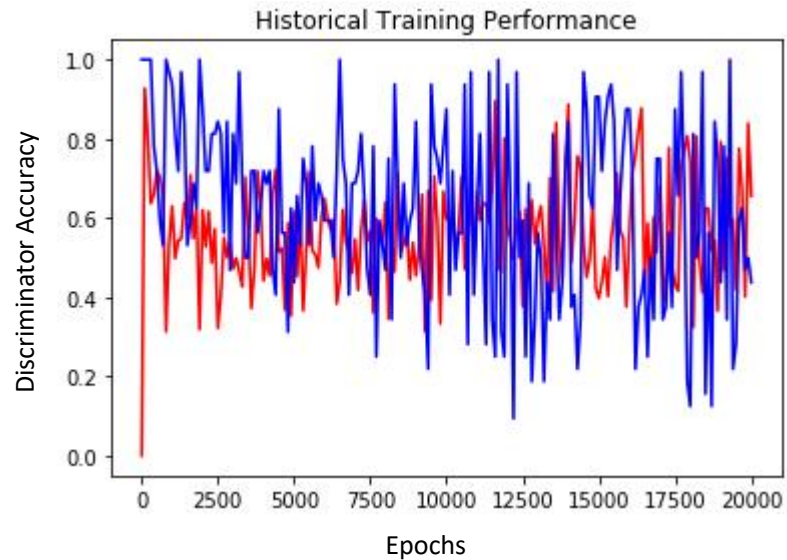
Epoch 1,000



Epoch 20,000



Adversarial game & performance



Deployment Performance

Normal Accuracy: 100%

Anomaly Accuracy: 95%

How can we improve upon the GAN framework?

✓	Standard	Vanilla GAN	Point-in-time observations, no context
✓ ✓	Good	RNN GAN	Sequence of observations, short-term context only
✓ ✓ ✓	Best	LSTM GAN	Sequence of observations, short and long-term context